

Lesson13 – Running Container with non-root User

In this lesson we will deploy MSSQL with persistent volume for high availability and non-root user.

The problem of running container with root user

1. Using shared volume mounted into several containers like central logging folder. The contained process running as root will have full access to every path on that mount volume. Now the application has a directory traversal vulnerability and hacker can poke around the container file system and find some config folder for "another application" on that directory. Even worse, if we add an hostpath volume, bind mount from container to the node VM like our MSMSQL container - now root process can read and write to any path on the host file system anywhere under that mount point.

2. hacker can setup a vulnerability with remote code execution in the container. When root user run the container it has a free access in the container file system to add and modify executable files, install packages and pretty much have a way in there. This raises a risk that some hacker will setup a vulnerability with remote code execution in the container.

Verify that root user run the MSSQL container:

```
kubectl create namespace employee
```

```
kubectl create secret generic mssql-secret --namespace=employee --from-literal='ConnectionString="server=mssql-service;Initial Catalog=EmployeeDB;Persist Security Info=False;User ID=sa;Password=MyDemoPwd2021!;MultipleActiveResultSets=true"' --from-literal='SA_PASSWORD=MyDemoPwd2021!'
```

```
cd C:\kubernetes\kubernetes-security\deployment
```

```
kubectl apply -f .\mssql-deploy-with-secret-and-pv.yml
```

```
kubectl get all -n employee
```

```
PS C:\Kubernetes\EmployeesManagement\manifests> kubectl get all -n employee
```

NAME	READY	STATUS	RESTARTS	AGE
pod/mssql-deployment-6bcb97764c-2tpl4	1/1	Running	0	4m18s

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
service/mssql-service	LoadBalancer	10.102.148.234	localhost	1433:30123/TCP	4m18s

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
deployment.apps/mssql-deployment	1/1	1	1	4m18s

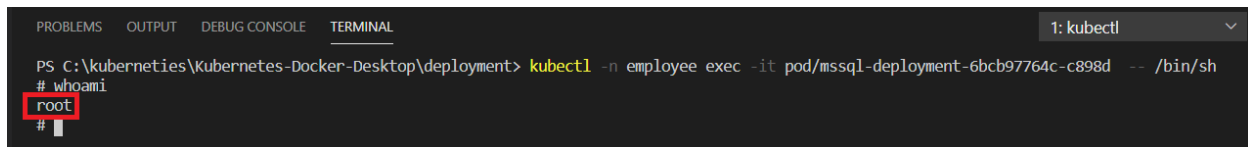
NAME	DESIRED	CURRENT	READY	AGE
replicaset.apps/mssql-deployment-6bcb97764c	1	1	1	4m18s

```
PS C:\Kubernetes\EmployeesManagement\manifests>
```

```
kubectl -n employee exec -it pod/mssql-deployment-6bcb97764c-2tpl4 -- /bin/sh
```

whoami

exit

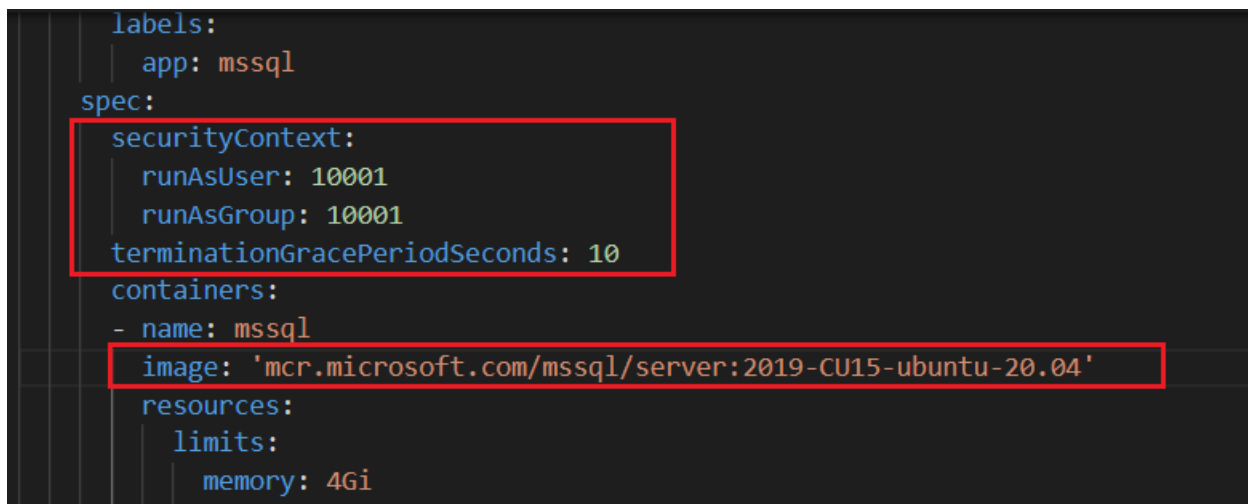
A terminal window with a dark background. The title bar shows '1: kubectl'. The command prompt is 'PS C:\kubernetes\Kubernetes-Docker-Desktop\deployment> kubectl -n employee exec -it pod/mssql-deployment-6bcb97764c-c898d -- /bin/sh'. The output shows a shell prompt '#', followed by 'root' (highlighted with a red box), and another shell prompt '#'.

```
PS C:\kubernetes\Kubernetes-Docker-Desktop\deployment> kubectl -n employee exec -it pod/mssql-deployment-6bcb97764c-c898d -- /bin/sh
# whoami
root
#
```

We can see that **root user** runs the MSSQL container.

Solution:

Create **mssql-deploy-with-secret-pv-non-root.yml** and add userid 10001 in the container layer as below:

A code editor showing a YAML file. The content is a Kubernetes deployment manifest. A red box highlights the 'securityContext' block, and another red box highlights the 'image' field in the 'containers' list.

```
labels:
  app: mssql
spec:
  securityContext:
    runAsUser: 10001
    runAsGroup: 10001
    terminationGracePeriodSeconds: 10
  containers:
    - name: mssql
      image: 'mcr.microsoft.com/mssql/server:2019-CU15-ubuntu-20.04'
      resources:
        limits:
          memory: 4Gi
```

kubectl delete ns employee

kubectl create namespace employee

kubectl create secret generic mssql-secret --namespace=employee --from-literal='ConnectionString="server=mssql-service;Initial Catalog=EmployeeDB;Persist Security Info=False;User ID=sa;Password=MyDemoPwd2021!;MultipleActiveResultSets=true"' --from-literal='SA_PASSWORD=MyDemoPwd2021!'

kubectl apply -f .\mssql-deploy-with-secret-pv-non-root.yml

kubectl get all -n employee

```
PS C:\kubernetes\kubernetes-security\deployment> kubectl get all -n employee
NAME                                READY   STATUS    RESTARTS   AGE
pod/mssql-deployment-5bc598bfd8-26pw7 1/1     Running   0           15m

NAME                                TYPE          CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
service/mssql-service               LoadBalancer  10.99.164.54  localhost     1433:31664/TCP   15m

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/mssql-deployment    1/1     1             1           15m

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/mssql-deployment-5bc598bfd8 1         1         1       15m
PS C:\kubernetes\kubernetes-security\deployment>
```

kubectl -n employee exec -it pod/mssql-deployment-5bc598bfd8-26pw7 -- /bin/sh

whoami

id

ps aux

```
PS C:\kubernetes\kubernetes-security\deployment> kubectl -n employee exec -it pod/mssql-deployment-54bd89c765-dqjkb -- /bin/sh
$ whoami
mssql
$ id
uid=10001(mssql) gid=10001 groups=10001
$ ps aux
USER          PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
mssql         1  0.1  0.3  61624 23608 ?        Ssl   16:52   0:00 /opt/mssql/bin/sqlservr
mssql        10 10.5 13.0 16071644 833188 ?        Sl    16:52   0:33 /opt/mssql/bin/sqlservr
mssql       265  0.0  0.0   2612    604 pts/0    Ss    16:57   0:00 /bin/sh
mssql       296  0.0  0.0   5900   2836 pts/0    R+    16:57   0:00 ps aux
$
```

top

q

```
top - 16:58:08 up 2:46, 0 users, load average: 0.43, 0.64, 0.54
Tasks: 4 total, 1 running, 3 sleeping, 0 stopped, 0 zombie
%Cpu(s): 2.5 us, 2.1 sy, 0.0 ni, 94.8 id, 0.0 wa, 0.0 hi, 0.6 si, 0.0 st
MiB Mem : 6237.4 total, 128.4 free, 2158.4 used, 3950.6 buff/cache
MiB Swap: 2048.0 total, 2045.2 free, 2.8 used. 3422.7 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
10	mssql	20	0	15.3g	848972	61784	S	7.7	13.3	0:36.50	sqlservr
1	mssql	20	0	61624	23608	9992	S	0.0	0.4	0:00.46	sqlservr
265	mssql	20	0	2612	604	540	S	0.0	0.0	0:00.01	sh
313	mssql	20	0	6144	3316	2800	R	0.0	0.1	0:00.00	top

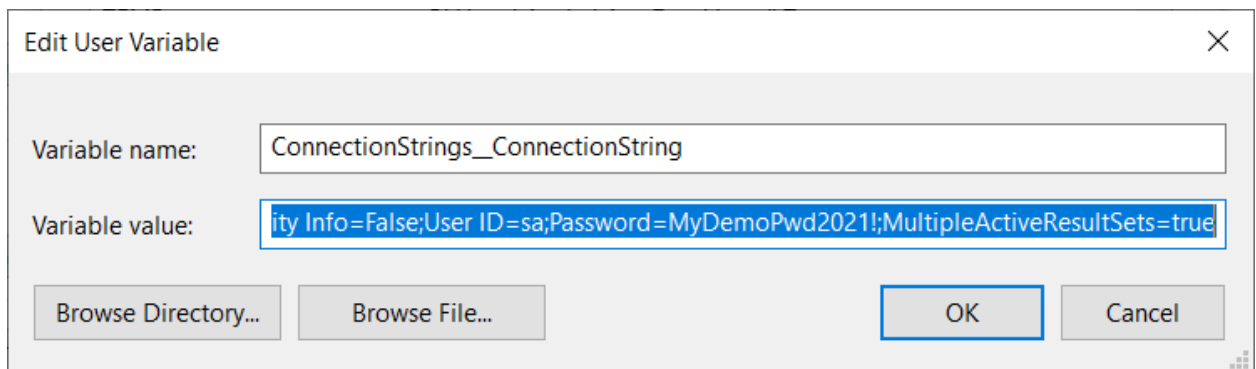
We can see that **non-root user** running the MSSQL container process with **user mssql (10001)**

```
cd C:\kubernetes\kubernetes-security\Employees
```

```
dotnet ef database update
```

NOTE: before running the command above make sure that the connection string exists in User Variable and close VScode and CMD and run again:

```
server=localhost,1433;Initial Catalog=EmployeeDB;Persist Security Info=False;User ID=sa;Password=MyDemoPwd2021!;MultipleActiveResultSets=true
```



Done!